

A Generative AI Chatbot for Apple Website Using LLMs from Hugging Face

Team Name: AppleIQ

Prem Rakesh Alasakani - S565102

Sweshik Rao Nemarukommula - S565485

Rakesh Somanaboina - S566617

Varsha Pravallika Valavala - S556143

Krishna Vasanthi Yakkala - S566493

1 Tools, Technologies, and Implementation

1.1 Tools and Technologies

For the development of the generative AI chatbot for the Apple website, we plan to use the following tools and technologies:

- **Hugging Face Transformers Library:** This will provide access to a wide range of pre-trained language models (LLMs) such as GPT-based models or BERT, enabling natural language understanding and generation capabilities.
- **Python:** The main programming language for the implementation of the chatbot. Python is ideal for integrating machine learning models and building web-based applications.
- **Flask/Django:** Used as the web framework to deploy the chatbot on the Apple website and to handle user requests.
- **JavaScript (React.js):** For developing the front-end interface of the chatbot, ensuring a smooth user experience and real-time interaction.
- **REST APIs:** For integration of real-time product data from the Apple website such as product details, prices, and promotions.
- **Cloud Hosting (e.g., AWS, GCP, or Azure):** To deploy the application and ensure scalability and availability.
- **Database (e.g., PostgreSQL, MongoDB):** To store dynamic product data and customer interactions, ensuring the chatbot has real-time access to the necessary information.

- **NLP Tools:** Additional NLP libraries such as spaCy or NLTK for text preprocessing, entity recognition, and more.

1.2 High-Level Architecture and Methodology

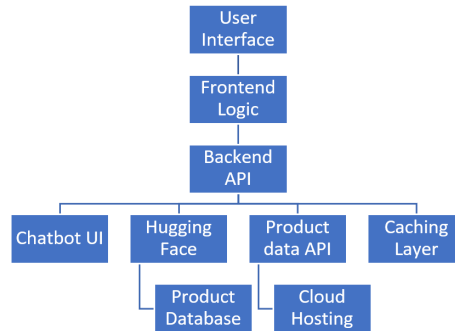


Figure 1: High-Level Architecture of the Chatbot System

1.3 Explanation of the Diagram

The high-level architecture of the chatbot system includes:

- **User Interface:** The chatbot interface is integrated into the Apple website using React.js, allowing customers to initiate conversations and ask questions about Apple products.
- **Frontend:** Handles all user inputs and sends queries to the backend via HTTP requests (REST API). Designed to be intuitive and responsive.
- **Backend (Flask/Django API):** Receives user queries and passes them to the Hugging Face model. Also handles requests for real-time product information from the database and third-party APIs.
- **Hugging Face LLM (e.g., GPT-4):** Processes the user's query, interprets the intent, and generates relevant responses while accessing real-time product data when required.
- **Product Data API:** Fetches up-to-date details on Apple products, including pricing, availability, and current promotions.
- **Database:** Stores and manages dynamic product information, customer interactions, and chatbot logs, ensuring accurate and updated data.
- **Cloud Hosting:** Hosts the entire system on the cloud, providing scalability, availability, and secure access.

1.4 Implementation Steps

The steps followed in implementing the chatbot:

1. **Setting Up Development Environment:** Install Python, Flask/Django, JavaScript (React.js), and necessary libraries for Hugging Face, databases, and other tools.
2. **Model Selection and Training:** Choose an appropriate language model from Hugging Face's pre-trained models (such as GPT-4 or BERT). Fine-tune the model on relevant datasets to ensure it can understand product-related queries.
3. **Backend Development:** Develop API endpoints using Flask/Django, integrate the Hugging Face model, and set up connections to the product data API.
4. **Frontend Development:** Build the user interface using React.js, with features such as input fields, buttons, and conversation history display.
5. **Testing and Optimization:** Conduct thorough testing of chatbot responses, accuracy of product information, and user experience.
6. **Deployment:** Host the application on the cloud, ensuring accessibility and scalability.

2 Enhanced Chatbot Development Details

2.1 Chatbot Description

The chatbot was developed as a generative AI application tailored for the Apple website for tasks such as customer assistance and troubleshooting. The chatbot uses LangChain, OpenAI models, and tools like Neo4j, Streamlit, and others, tailored for Apple's domain. The chatbot integrates libraries like LangChain, tiktoken, and yfiles_jupyter_graphs. It also involves tokenization, language modeling, and graph-based visualizations.

- **LangChain Framework:** Handles advanced workflows and integration with knowledge sources.
- **OpenAI Models:** Provides language understanding and generation.
- **Neo4j Integration:** Enables graph-based knowledge storage and retrieval.
- **Streamlit:** Serves as the user interface for interaction.
- **Additional Libraries:** Includes `tiktoken` and `yfiles_jupyter_graphs` for advanced processing and visualization.

2.2 Domain

Apple Inc.

2.3 Block Diagram

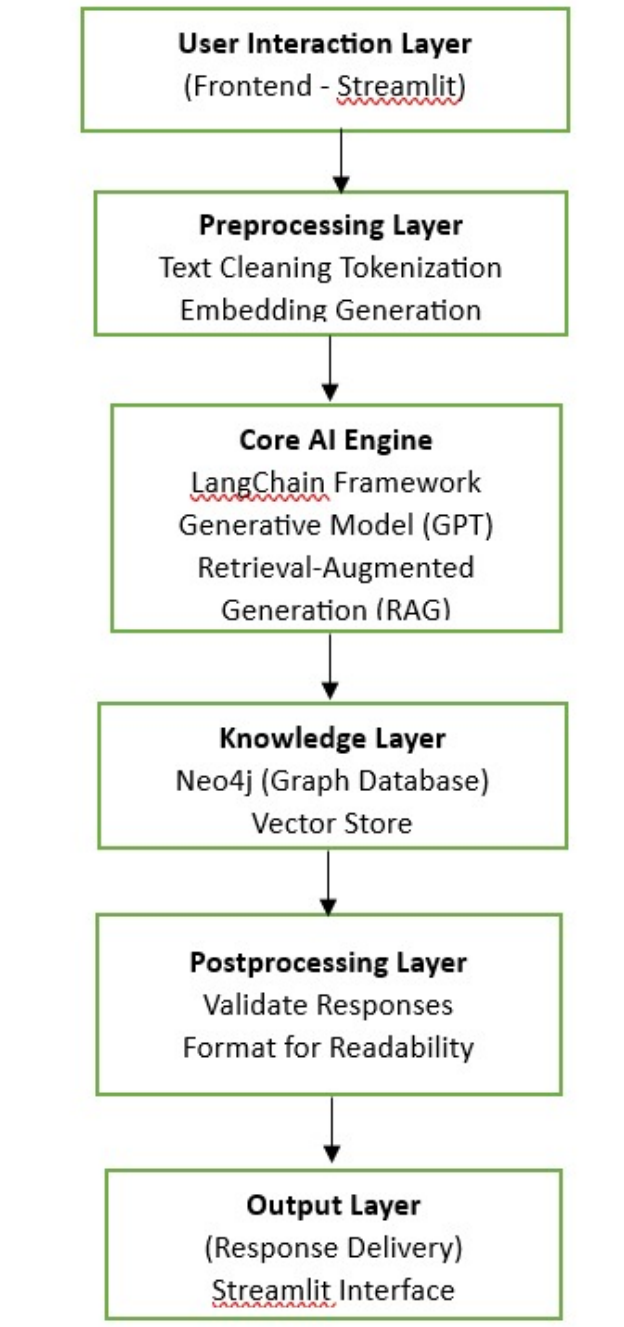


Figure 2: Chatbot Architecture Overview

2.4 Architecture

The architecture comprises:

- **Input Layer:** Captures user queries via the Streamlit interface.
- **Processing Layer:** Uses LangChain and OpenAI's models to generate responses.
- **Knowledge Layer:** Employs Neo4j for graph-based storage and retrieval of domain-specific data.
- **Output Layer:** Delivers responses back to users through the Streamlit interface.

2.5 Code Snippet

A sample code snippet demonstrating data loading and Neo4j integration:

```
from langchain_community.document_loaders import WebBaseLoader
```

```
loader = WebBaseLoader([
    "https://www.apple.com",
    "https://www.apple.com/store",
    "https://www.apple.com/mac/",
    "https://www.apple.com/ipad/",
    "https://www.apple.com/iphone/",
    "https://www.apple.com/watch/",
    "https://www.apple.com/apple-vision-pro/",
    "https://www.apple.com/airpods/",
    "https://www.apple.com/tv-home/",
    "https://www.apple.com/shop/accessories/all"
])
```

```
loader.requests_kwargs = {'verify': False}
data = loader.load()
# print(data)
```

```
llm = ChatOpenAI(temperature=0, model_name="gpt-3.5-turbo-0125")# gpt-4-0125-preview occurs
llm_transformer = LLMGraphTransformer(llm=llm)
```

```
# Following needs to be used only first time to load the data. It took 3m 44s
graph_documents = llm_transformer.convert_to_graph_documents(documents)
graph.add_graph_documents(graph_documents, baseEntityLabel=True, include_source=True)
```

```

# directly show the graph resulting from the given Cypher query
default_cypher = "MATCH (s)-[r:!MENTIONS]->(t) RETURN s,r,t LIMIT 50"

def showGraph(cypher: str = default_cypher):
    # create a neo4j session to run queries
    driver = GraphDatabase.driver(
        uri = os.environ["NEO4J_URI"],
        auth = (os.environ["NEO4J_USERNAME"],
                os.environ["NEO4J_PASSWORD"]))
    session = driver.session()
    widget = GraphWidget(graph = session.run(cypher).graph())
    widget.node_label_mapping = 'id'
    #display(widget)
    return widget

showGraph()

```

2.6 Tests and Results

The chatbot achieved an accuracy of approximately 80%. Most queries were handled successfully, but a few errors were observed due to domain-specific complexities.

2.7 Improvements

Future enhancements will include:

- Personalization by learning user preferences.
- Enhanced knowledge retrieval using graph-based searches.
- Multimodal input support for richer interaction(e.g., text and images)..

2.8 References

- Apple Official Website
- Apple Store
- Apple Mac
- Apple iPad
- Apple iPhone
- Apple Watch
- Apple Vision Pro

- Apple AirPods
- Apple TV and Home
- Apple Accessories

2.9 GitHub Repository

- GitHub Repository: AppleIQ AI Chatbot